

# Parallel External Memory Suffix Sorting

Juha Kärkkäinen   Dominik Kempa   Simon J. Puglisi

University of Helsinki, Finland

CPM, Ischia, July 2015

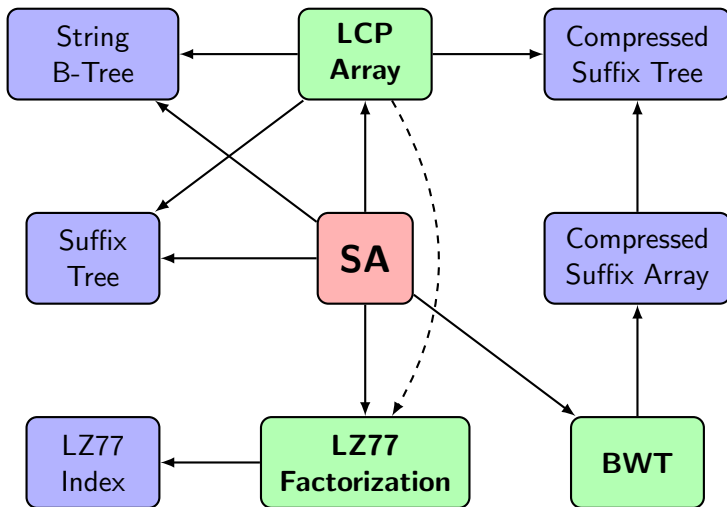
# Outline

- 1 Introduction and Background
- 2 Internal memory SACA
- 3 External memory SACA

# Outline

- 1 Introduction and Background
- 2 Internal memory SACA
- 3 External memory SACA

## Central Role of Suffix Array (SA)



# Suffix Array Construction

Suffix array construction algorithm (SACA):

Sort all suffixes of a text  
lexicographically

SA

6	\$
5	a\$
3	ana\$
1	anana\$
0	banana\$
4	na\$
2	nana\$

	Internal memory	External memory
Sequential	<ul style="list-style-type: none"> <li>Manber, Myers (1993)</li> <li>Sadakane (1998)</li> <li>Itoh, Tanaka (1999)</li> <li>Larsson, Sadakane (1999)</li> <li>Manzini, Ferragina (2002,2004)</li> <li>Burkhardt, Kärkkäinen (2003)</li> <li>Kärkkäinen, Sanders, Burkhardt (2003,2006)</li> <li>Kim, Sim, Park, Park (2003,2005)</li> <li>Ko, Aluru (2003,2005)</li> <li>Hon, Sadakane, Sung (2003,2009)</li> <li>Schürmann, Stoye (2005,2007)</li> <li>Maniscalco, Puglisi (2007)</li> <li>divsufsort: Mori (2008)</li> <li>Nong, Zhang, Chan (2011)</li> <li>Nong (2013)</li> </ul>	<ul style="list-style-type: none"> <li>Gonnet, Baeza-Yates, Snider (1987,1992)</li> <li>Crauser, Ferragina (1999,2002)</li> <li>Kärkkäinen, Sanders, Burkhardt (2003,2006)</li> <li>Dementiev, Kärkkäinen, Mehnert, Sanders (2005,2008)</li> <li>Ferragina, Gagie, Manzini (2010,2012)</li> <li>Beller, Zwerger, Gog, Ohlebusch (2013)</li> <li>eSAIS: Bingmann, Fischer, Osipov (2013)</li> <li>SAScan: Kärkkäinen, K (2014)</li> <li>Tischler (2014)</li> <li>Nong, Chan, Zhang, Guan (2014)</li> <li>Nong, Chan, Hu, Wu (2015)</li> </ul>
Parallel	<ul style="list-style-type: none"> <li>Kulla, Sanders (2006,2007)</li> <li>pDC3: Blelloch, Shun (2011)</li> <li>Osipov (2012)</li> <li>Deo, Keely (2013)</li> <li><a href="#">pSAScan: This paper</a></li> </ul>	<ul style="list-style-type: none"> <li><a href="#">pSAScan: This paper</a></li> </ul>

# Outline

- 1 Introduction and Background
- 2 Internal memory SACA**
- 3 External memory SACA

	Internal memory	External memory
Sequential	<ul style="list-style-type: none"> <li>● Manber, Myers (1993)</li> <li>● Sadakane (1998)</li> <li>● Itoh, Tanaka (1999)</li> <li>● Larsson, Sadakane (1999)</li> <li>● Manzini, Ferragina (2002,2004)</li> <li>● Burkhardt, Kärkkäinen (2003)</li> <li>● Kärkkäinen, Sanders, Burkhardt (2003,2006)</li> <li>● Kim, Sim, Park, Park (2003,2005)</li> <li>● Ko, Aluru (2003,2005)</li> <li>● Hon, Sadakane, Sung (2003,2009)</li> <li>● Schürmann, Stoye (2005,2007)</li> <li>● Maniscalco, Puglisi (2007)</li> <li>● <b>divsufsort: Mori (2008)</b></li> <li>● Nong, Zhang, Chan (2011)</li> <li>● Nong (2013)</li> </ul>	<ul style="list-style-type: none"> <li>● Gonnet, Baeza-Yates, Snider (1987,1992)</li> <li>● Crauser, Ferragina (1999,2002)</li> <li>● Kärkkäinen, Sanders, Burkhardt (2003,2006)</li> <li>● Dementiev, Kärkkäinen, Mehnert, Sanders (2005,2008)</li> <li>● Ferragina, Gagie, Manzini (2010,2012)</li> <li>● Beller, Zwerger, Gog, Ohlebusch (2013)</li> <li>● <b>eSAIS: Bingmann, Fischer, Osipov (2013)</b></li> <li>● <b>SAscan: Kärkkäinen, K (2014)</b></li> <li>● Tischler (2014)</li> <li>● Nong, Chan, Zhang, Guan (2014)</li> <li>● Nong, Chan, Hu, Wu (2015)</li> </ul>
Parallel	<ul style="list-style-type: none"> <li>● Kulla, Sanders (2006,2007)</li> <li>● <b>pDC3: Blelloch, Shun (2011)</b></li> <li>● Osipov (2012)</li> <li>● Deo, Keely (2013)</li> <li>● <b>pSAscan: This paper</b></li> </ul>	<ul style="list-style-type: none"> <li>● <b>pSAscan: This paper</b></li> </ul>



# divsufsort

- Currently, the fastest **sequential** internal memory SACA
- Time complexity:  $\mathcal{O}(n \log n)$
- Available 32- and 64-bit versions using  **$5n$**  and  **$9n$**  bytes of RAM

# divsufsort

- Currently, the fastest **sequential** internal memory SACA
- Time complexity:  $\mathcal{O}(n \log n)$
- Available 32- and 64-bit versions using  **$5n$**  and  **$9n$**  bytes of RAM

# divsufsort

- Currently, the fastest **sequential** internal memory SACA
- Time complexity:  $\mathcal{O}(n \log n)$
- Available 32- and 64-bit versions using  **$5n$**  and  **$9n$**  bytes of RAM

# pDC3

- Parallel version of DC3 algorithm of Kärkkäinen, Sanders, Burkhardt (2003,2006)
- Fastest parallel SACA multi-core architecture
- Optimal  $\mathcal{O}(n)$  work
- Implementation by Blelloch and Shun
- RAM usage:  $21n$  (32-bit),  $41n$  (64-bit)
- (divsufsort:  $5n$  (32-bit),  $9n$  (64-bit))

# pDC3

- Parallel version of DC3 algorithm of Kärkkäinen, Sanders, Burkhardt (2003,2006)
- Fastest parallel SACA multi-core architecture
- Optimal  $\mathcal{O}(n)$  work
- Implementation by Blelloch and Shun
- RAM usage:  $21n$  (32-bit),  $41n$  (64-bit)
- (divsufsort:  $5n$  (32-bit),  $9n$  (64-bit))

# pDC3

- Parallel version of DC3 algorithm of Kärkkäinen, Sanders, Burkhardt (2003,2006)
- Fastest parallel SACA multi-core architecture
- Optimal  $\mathcal{O}(n)$  work
- Implementation by Blelloch and Shun
- RAM usage:  $21n$  (32-bit),  $41n$  (64-bit)
- (divsufsort:  $5n$  (32-bit),  $9n$  (64-bit))

# pDC3

- Parallel version of DC3 algorithm of Kärkkäinen, Sanders, Burkhardt (2003,2006)
- Fastest parallel SACA multi-core architecture
- Optimal  $\mathcal{O}(n)$  work
- Implementation by Blelloch and Shun
- RAM usage:  $21n$  (32-bit),  $41n$  (64-bit)
- (divsufsort:  $5n$  (32-bit),  $9n$  (64-bit))

## pDC3

- Parallel version of DC3 algorithm of Kärkkäinen, Sanders, Burkhardt (2003,2006)
- Fastest parallel SACA multi-core architecture
- Optimal  $\mathcal{O}(n)$  work
- Implementation by Blelloch and Shun
- RAM usage:  $21n$  (32-bit),  $41n$  (64-bit)
- (divsufsort:  $5n$  (32-bit),  $9n$  (64-bit))



# SAscan: general overview

- 1 Divide text into blocks
- 2 Construct SA for each block
- 3 Merge each block SA into full SA

# SAscan: general overview


- 1 Divide text into blocks
- 2 Construct SA for each block
- 3 Merge each block SA into full SA

# SAscan: general overview


- 1 Divide text into blocks
- 2 Construct SA for each block
- 3 Merge each block SA into full SA

# SAscan: general overview

- 1 Divide text into blocks
- 2 Construct SA for each block
- 3 Merge each block SA into full SA
  - multiple merge schedules possible



Unbalanced merge



Fully balanced merge

## SAscan: parallel in-RAM version

- 1 Divide text into blocks
- 2 Construct SA for each block *for all blocks in parallel*
- 3 Merge each block SA into full SA

## SAscan: parallel in-RAM version


- 1 Divide text into blocks
- 2 Construct SA for each block **for all blocks in parallel**
- 3 Merge each block SA into full SA

## SAscan: parallel in-RAM version

- 1 Divide text into blocks
- 2 Construct SA for each block **for all blocks in parallel**
- 3 Merge each block SA into full SA

# SAscan: parallel in-RAM version


- 1 Divide text into blocks
- 2 Construct SA for each block **for all blocks in parallel**
- 3 Merge each block SA into full SA
  - **balanced schedule**
  - **how to parallelize?**






# SAscan: parallel in-RAM version


- 1 Divide text into blocks
- 2 Construct SA for each block **for all blocks in parallel**
- 3 Merge each block SA into full SA
  - **balanced schedule**
  - **how to parallelize?**




## Merging SA of two blocks




## Merging SA of two blocks



## Merging SA of two blocks



## Merging SA of two blocks



## Merging SA of two blocks: parallel in-RAM version

### ① $SA_Y + Z \rightarrow \text{gap}_Y$

- backward search with  $Y$  as text and  $Z$  as pattern [Ferragina, Gagie & Manzini (2010,2012)]
- key property: needs rank data structure on BWT of  $Y$
- **parallelization**: start backward search from multiple positions

### ② $SA_Y + SA_Z + \text{gap}_Y \rightarrow SA_{YZ}$

- **parallelization**: possible but complex, details omitted

## Merging SA of two blocks: parallel in-RAM version

- ①  $SA_Y + Z \rightarrow \text{gap}_Y$ 
  - backward search with  $Y$  as text and  $Z$  as pattern [Ferragina, Gaggie & Manzini (2010,2012)]
  - key property: needs rank data structure on BWT of  $Y$
  - **parallelization**: start backward search from multiple positions
- ②  $SA_Y + SA_Z + \text{gap}_Y \rightarrow SA_{YZ}$ 
  - **parallelization**: possible but complex, details omitted

## Merging SA of two blocks: parallel in-RAM version

- ①  $SA_Y + Z \rightarrow \text{gap}_Y$ 
  - backward search with  $Y$  as text and  $Z$  as pattern [Ferragina, Gagie & Manzini (2010,2012)]
  - key property: needs rank data structure on BWT of  $Y$
  - **parallelization**: start backward search from multiple positions
- ②  $SA_Y + SA_Z + \text{gap}_Y \rightarrow SA_{YZ}$ 
  - **parallelization**: possible but complex, details omitted



## Merging SA of two blocks: parallel in-RAM version

- ①  $SA_Y + Z \rightarrow \text{gap}_Y$ 
  - backward search with  $Y$  as text and  $Z$  as pattern  
[Ferragina, Gagie & Manzini (2010,2012)]
  - key property: needs rank data structure on BWT of  $Y$
  - **parallelization**: start backward search from multiple positions
- ②  $SA_Y + SA_Z + \text{gap}_Y \rightarrow SA_{YZ}$ 
  - **parallelization**: possible but complex, details omitted

## Merging SA of two blocks: parallel in-RAM version

- ①  $SA_Y + Z \rightarrow \text{gap}_Y$ 
  - backward search with  $Y$  as text and  $Z$  as pattern  
[Ferragina, Gaggie & Manzini (2010,2012)]
  - key property: needs rank data structure on BWT of  $Y$
  - **parallelization**: start backward search from multiple positions
- ②  $SA_Y + SA_Z + \text{gap}_Y \rightarrow SA_{YZ}$ 
  - **parallelization**: possible but complex, details omitted

## Merging SA of two blocks: parallel in-RAM version

- ①  $SA_Y + Z \rightarrow \text{gap}_Y$ 
  - backward search with  $Y$  as text and  $Z$  as pattern  
[Ferragina, Gagie & Manzini (2010,2012)]
  - key property: needs rank data structure on BWT of  $Y$
  - **parallelization**: start backward search from multiple positions
- ②  $SA_Y + SA_Z + \text{gap}_Y \rightarrow SA_{YZ}$ 
  - **parallelization**: possible but complex, details omitted

# Internal memory pSAscan - summary

- Work:  $\mathcal{O}(n \log p)$
- Space usage
  - $9n$  (32-bit),  $10n$  (40-bit)
  - (pDC3:  $21n$  (32-bit),  $41n$  (64-bit))

## Internal memory pSAscan - summary

- Work:  $\mathcal{O}(n \log p)$
- Space usage
  - $9n$  (32-bit),  $10n$  (40-bit)
  - (pDC3:  $21n$  (32-bit),  $41n$  (64-bit))

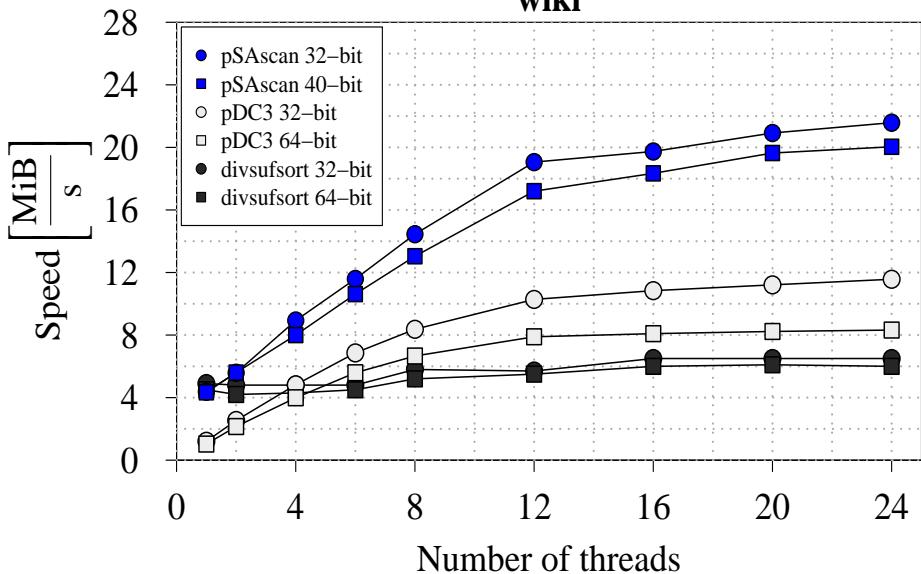
## Internal memory pSAscan - summary

- Work:  $\mathcal{O}(n \log p)$
- Space usage
  - $9n$  (32-bit),  $10n$  (40-bit)
  - (pDC3:  $21n$  (32-bit),  $41n$  (64-bit))

# Experiments

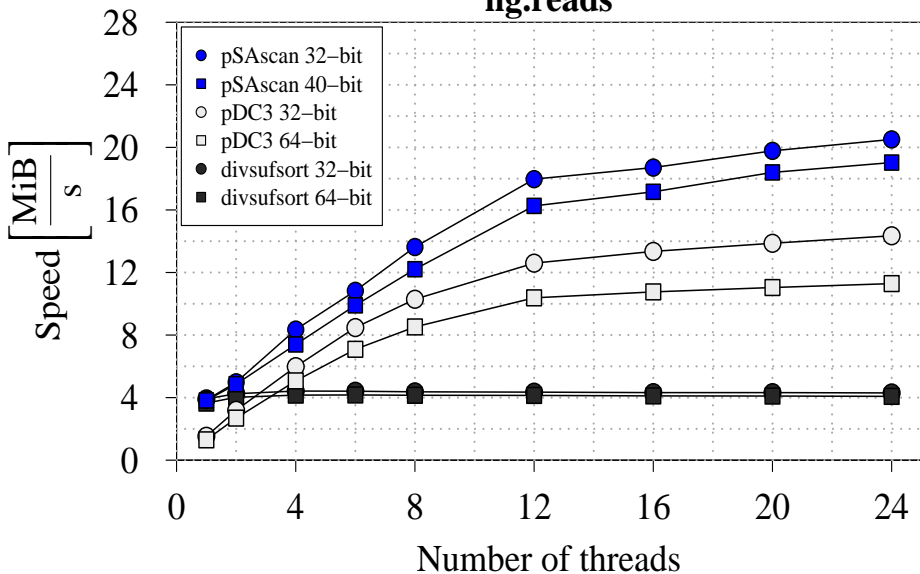
(comparison of internal memory parallel  
suffix array construction algorithms)

wiki

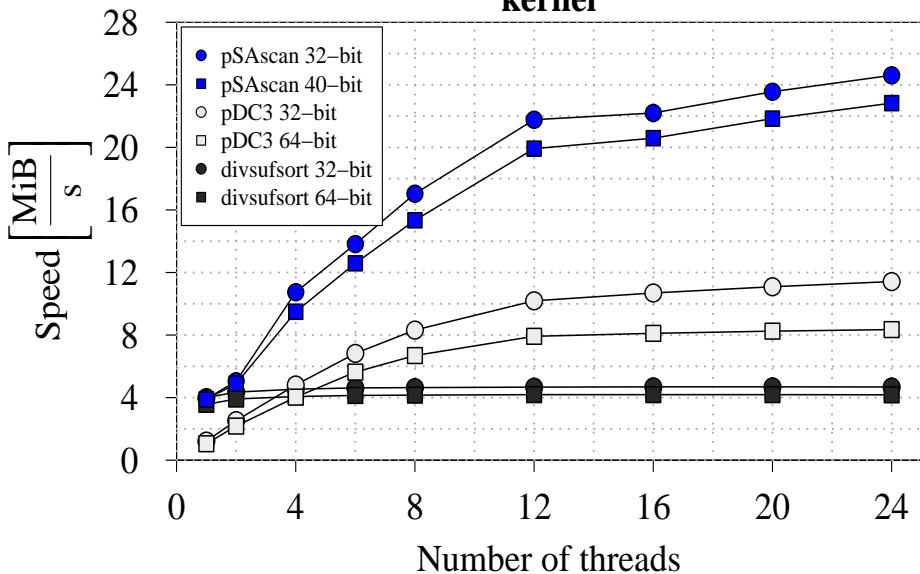





# hg.reads



# kernel



# random (sigma = 256)



# Outline

- 1 Introduction and Background
- 2 Internal memory SACA
- 3 External memory SACA

	Internal memory	External memory
Sequential	<ul style="list-style-type: none"> <li>Manber, Myers (1993)</li> <li>Sadakane (1998)</li> <li>Itoh, Tanaka (1999)</li> <li>Larsson, Sadakane (1999)</li> <li>Manzini, Ferragina (2002,2004)</li> <li>Burkhardt, Kärkkäinen (2003)</li> <li>Kärkkäinen, Sanders, Burkhardt (2003,2006)</li> <li>Kim, Sim, Park, Park (2003,2005)</li> <li>Ko, Aluru (2003,2005)</li> <li>Hon, Sadakane, Sung (2003,2009)</li> <li>Schürmann, Stoye (2005,2007)</li> <li>Maniscalco, Puglisi (2007)</li> <li>divsufsort: Mori (2008)</li> <li>Nong, Zhang, Chan (2011)</li> <li>Nong (2013)</li> </ul>	<ul style="list-style-type: none"> <li>Gonnet, Baeza-Yates, Snider (1987,1992)</li> <li>Crauser, Ferragina (1999,2002)</li> <li>Kärkkäinen, Sanders, Burkhardt (2003,2006)</li> <li>Dementiev, Kärkkäinen, Mehnert, Sanders (2005,2008)</li> <li>Ferragina, Gagie, Manzini (2010,2012)</li> <li>Beller, Zwerger, Gog, Ohlebusch (2013)</li> <li>eSAIS: Bingmann, Fischer, Osipov (2013)</li> <li>SAScan: Kärkkäinen, K (2014)</li> <li>Tischler (2014)</li> <li>Nong, Chan, Zhang, Guan (2014)</li> <li>Nong, Chan, Hu, Wu (2015)</li> </ul>
Parallel	<ul style="list-style-type: none"> <li>Kulla, Sanders (2006,2007)</li> <li>pDC3: Blelloch, Shun (2011)</li> <li>Osipov (2012)</li> <li>Deo, Keely (2013)</li> <li>pSAscan: This paper</li> </ul>	<ul style="list-style-type: none"> <li>pSAscan: This paper</li> </ul>

- Optimal  $\mathcal{O}(\text{sort}(n))$  I/O complexity
  - $\text{sort}(n)$  = complexity of sorting  $n$  integers
  - $\mathcal{O}\left(\frac{n}{B} \log_{M/B} \frac{n}{B}\right)$  I/Os
  - $M$  = size of RAM,  $B$  = size of disk block (in units of  $\log n$  bits)
- Implementation by Bingmann, Fischer & Osipov (2013)
- $28n$  bytes of disk space
  - $n$  bytes for input
  - $5n$  bytes output

- Optimal  $\mathcal{O}(\text{sort}(n))$  I/O complexity
  - $\text{sort}(n)$  = complexity of sorting  $n$  integers
  - $\mathcal{O}\left(\frac{n}{B} \log_{M/B} \frac{n}{B}\right)$  I/Os
  - $M$  = size of RAM,  $B$  = size of disk block (in units of  $\log n$  bits)
- Implementation by Bingmann, Fischer & Osipov (2013)
- $28n$  bytes of disk space
  - $n$  bytes for input
  - $5n$  bytes output


- Optimal  $\mathcal{O}(\text{sort}(n))$  I/O complexity
  - $\text{sort}(n)$  = complexity of sorting  $n$  integers
  - $\mathcal{O}\left(\frac{n}{B} \log_{M/B} \frac{n}{B}\right)$  I/Os
  - $M$  = size of RAM,  $B$  = size of disk block  
(in units of  $\log n$  bits)
- Implementation by Bingmann, Fischer & Osipov (2013)
  - $28n$  bytes of disk space
    - $n$  bytes for input
    - $5n$  bytes output



- Optimal  $\mathcal{O}(\text{sort}(n))$  I/O complexity
  - $\text{sort}(n)$  = complexity of sorting  $n$  integers
  - $\mathcal{O}\left(\frac{n}{B} \log_{M/B} \frac{n}{B}\right)$  I/Os
  - $M$  = size of RAM,  $B$  = size of disk block  
(in units of  $\log n$  bits)
- Implementation by Bingmann, Fischer & Osipov (2013)
- $28n$  bytes of disk space
  - $n$  bytes for input
  - $5n$  bytes output


# SAscan: EM version

- ❶ Divide text into blocks
  - small enough for RAM processing
- ❷ Construct SA for each block in RAM
- ❸ Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
⇒ unbalanced merging




# SAscan: EM version

- ❶ Divide text into blocks
  - small enough for RAM processing
- ❷ Construct SA for each block in RAM
- ❸ Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
⇒ unbalanced merging




# SAscan: EM version

- 1 Divide text into blocks
  - small enough for RAM processing
- 2 Construct SA for each block in RAM
- 3 Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
 $\Rightarrow$  unbalanced merging




# SAscan: EM version

- 1 Divide text into blocks
  - small enough for RAM processing
- 2 Construct SA for each block in RAM
- 3 Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-side block  
 $\Rightarrow$  unbalanced merging



# SAscan: EM version

- ❶ Divide text into blocks
  - small enough for RAM processing
- ❷ Construct SA for each block in RAM
- ❸ Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
 $\Rightarrow$  unbalanced merging



## SAscan: EM version

- $\Theta(n/M)$  blocks
- Time and I/O proportional to  $n^2/M$
- Fastest SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 5$$

- Lightweight w.r.t. disk space, uses  $7.5n$  bytes
  - eSAIS:  $28n$

## SAscan: EM version

- $\Theta(n/M)$  blocks
- Time and I/O proportional to  $n^2/M$
- Fastest SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 5$$

- Lightweight w.r.t. disk space, uses  $7.5n$  bytes
  - eSAIS:  $28n$



## SAscan: EM version

- $\Theta(n/M)$  blocks
- Time and I/O proportional to  $n^2/M$
- Fastest SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 5$$

- Lightweight w.r.t. disk space, uses  $7.5n$  bytes
  - eSAIS:  $28n$

## SAscan: EM version


- $\Theta(n/M)$  blocks
- Time and I/O proportional to  $n^2/M$
- Fastest SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 5$$

- Lightweight w.r.t. disk space, uses  $7.5n$  bytes
  - eSAIS:  $28n$


# SAscan: parallel EM version

- 1 Divide text into blocks
  - small enough for RAM processing
- 2 Construct SA for each block in RAM
  - use internal memory pSAscan
- 3 Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
⇒ unbalanced merging
  - parallelization: start backward search from multiple positions




# SAscan: parallel EM version

- 1 Divide text into blocks
  - small enough for RAM processing
- 2 Construct SA for each block in RAM
  - use internal memory pSAscan
- 3 Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
⇒ unbalanced merging
  - parallelization: start backward search from multiple positions




# SAscan: parallel EM version

- 1 Divide text into blocks
  - small enough for RAM processing
- 2 Construct SA for each block in RAM
  - use internal memory pSAscan
- 3 Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
⇒ unbalanced merging
  - parallelization: start backward search from multiple positions




# SAscan: parallel EM version

- 1 Divide text into blocks
  - small enough for RAM processing
- 2 Construct SA for each block in RAM
  - use internal memory pSAscan
- 3 Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
⇒ unbalanced merging
  - parallelization: start backward search from multiple positions




# SAscan: parallel EM version

- 1 Divide text into blocks
  - small enough for RAM processing
- 2 Construct SA for each block in RAM
  - use internal memory pSAscan
- 3 Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
⇒ unbalanced merging
  - parallelization: start backward search from multiple positions



# SAscan: parallel EM version

- ❶ Divide text into blocks
  - small enough for RAM processing
- ❷ Construct SA for each block in RAM
  - use internal memory pSAscan
- ❸ Merge each block SA into full SA
  - key property: backward search needs rank data structure on BWT of left hand-size block  
⇒ unbalanced merging
  - parallelization: start backward search from multiple positions





## SAscan: parallel EM version

- $\Theta(n/M)$  blocks
- Time and I/O proportional to  $n^2/M$
- Fastest SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

(on 12-core machine)

- Lightweight w.r.t. disk space, uses  $7.5n$  bytes
  - eSAIS:  $28n$

## SAscan: parallel EM version

- $\Theta(n/M)$  blocks
- Time and I/O proportional to  $n^2/M$
- Fastest SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

(on 12-core machine)

- Lightweight w.r.t. disk space, uses  $7.5n$  bytes
  - eSAIS:  $28n$

## SAscan: parallel EM version

- $\Theta(n/M)$  blocks
- Time and I/O proportional to  $n^2/M$
- Fastest SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

(on 12-core machine)

- Lightweight w.r.t. disk space, uses  $7.5n$  bytes
  - eSAIS:  $28n$

## SAscan: parallel EM version

- $\Theta(n/M)$  blocks
- Time and I/O proportional to  $n^2/M$
- Fastest SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$


(on 12-core machine)

- Lightweight w.r.t. disk space, uses  $7.5n$  bytes
  - eSAIS:  $28n$


# Experiments

(comparison of EM suffix array  
construction algorithms)

# kernel



# hg.reads



- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB



- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB

- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB

- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB

- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB

- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB

- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB

- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB

- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB



- 200GiB prefix of hg.reads

- Runtime

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	8.3 days	4.1 days
pSAscan	7.0 days	0.5 days

- Peak disk space usage

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	4.6 TiB	4.6 TiB
pSAscan	1.4 TiB	1.4 TiB

- I/O volume

	3.5 GiB of RAM	120 GiB of RAM
eSAIS	52.0 TiB	36.1 TiB
pSAscan	43.8 TiB	4.9 TiB

- Full instance (1TiB) of hg.reads, using 120GiB of RAM

	Runtime	Peak disk usage	I/O volume
pSAscan	8.1 days	7.3 TiB	48.3 TiB

## Concluding Remarks

- New internal-memory parallel SACA
  - $2 \times$  faster and  $\sim 2.5 \times$  less RAM than pDC3
- First external-memory parallel SACA
  - Fastest EM SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

- uses  $\sim 4$  times less disk space than eSAIS

Possible further improvements

- smaller rank data structure
- distributed / GPU implementation?

## Concluding Remarks

- New internal-memory parallel SACA
  - $2 \times$  faster and  $\sim 2.5 \times$  less RAM than pDC3
- First external-memory parallel SACA
  - Fastest EM SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

- uses  $\sim 4$  times less disk space than eSAIS

Possible further improvements

- smaller rank data structure
- distributed / GPU implementation?

## Concluding Remarks

- New internal-memory parallel SACA
  - $2 \times$  faster and  $\sim 2.5 \times$  less RAM than pDC3
- First external-memory parallel SACA
  - Fastest EM SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

- uses  $\sim 4$  times less disk space than eSAIS

Possible further improvements

- smaller rank data structure
- distributed / GPU implementation?

## Concluding Remarks

- New internal-memory parallel SACA
  - $2 \times$  faster and  $\sim 2.5 \times$  less RAM than pDC3
- First external-memory parallel SACA
  - Fastest EM SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

- uses  $\sim 4$  times less disk space than eSAIS

Possible further improvements

- smaller rank data structure
- distributed / GPU implementation?

## Concluding Remarks

- New internal-memory parallel SACA
  - $2 \times$  faster and  $\sim 2.5 \times$  less RAM than pDC3
- First external-memory parallel SACA
  - Fastest EM SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

- uses  $\sim 4$  times less disk space than eSAIS

Possible further improvements

- smaller rank data structure
- distributed / GPU implementation?

## Concluding Remarks

- New internal-memory parallel SACA
  - $2 \times$  faster and  $\sim 2.5 \times$  less RAM than pDC3
- First external-memory parallel SACA
  - Fastest EM SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

- uses  $\sim 4$  times less disk space than eSAIS

Possible further improvements

- smaller rank data structure
- distributed / GPU implementation?

## Concluding Remarks

- New internal-memory parallel SACA
  - $2 \times$  faster and  $\sim 2.5 \times$  less RAM than pDC3
- First external-memory parallel SACA
  - Fastest EM SACA when

$$\frac{\text{text size}}{\text{RAM size}} < 80$$

- uses  $\sim 4$  times less disk space than eSAIS

Possible further improvements

- smaller rank data structure
- distributed / GPU implementation?



# Thank you!

Code:

[www.cs.helsinki.fi/group/pads/](http://www.cs.helsinki.fi/group/pads/)